

# Extracting Music Features

April 20, 2022

## 1 Namaste! So, what is our project about?

### 1.0.1 Musical Genre Classification of Audio Signals.

- Citation: Tzanetakis, George, and Perry Cook. "Musical genre classification of audio signals," IEEE Transactions on speech and audio processing, 2002.
- Link: <https://www.cs.cmu.edu/~gtzan/work/pubs/tsap02gtzan.pdf>

### 1.0.2 Import Libraries

```
[1]: %matplotlib inline
import librosa
import pandas as pd
import numpy as np
import librosa.display
import sklearn
import matplotlib.pyplot as plt
import IPython.display as ipd
import os, glob
```

### 1.0.3 Configurations

```
[2]: config = {
    "dataset": {
#         "input": "test-dataset",
#         "output": "test-data"
        "input": "music-genres-dataset",
        "output": "music-data"
    },
    "counter": 0,
    "max_counter": 5,
    "persist": False
}
```

#### 1.0.4 Load all music files and extract their IDs and genres

```
[3]: music_data = dict()
for filepath in glob.iglob('data/' + config["dataset"]["input"] + '/*',
↪ recursive=True):
    if os.path.isfile(filepath):
        # Check if it is a music file
        if filepath.endswith(".wav"):
            music_data[filepath] = dict()
            music_data[filepath]["filepath"] = filepath
            music_data[filepath]["filename"] = filepath.rsplit('/',1)[1]
            music_data[filepath]["genre"] = "UNKNOWN" if "test" in
↪ config["dataset"]["input"] else music_data[filepath]["filename"].split(".")
↪ ,1)[0]

            config["counter"] += 1
            if config["counter"] == config["max_counter"]:
                break
```

#### 1.0.5 No. of Music files

```
[4]: len(music_data)
```

```
[4]: 5
```

#### 1.0.6 It has the following Genres:

- Pop
- Blues
- Classical
- Country
- Disco
- Hiphop
- Jazz
- Metal
- Pop
- Reggae
- Rock

### 1.0.7 Let's analyze one SAMPLE\_FILE and explore the features it provides..

```
[5]: SAMPLE_FILE = list(music_data.keys())[0]
SAMPLE_FILE
```

```
[5]: 'data/music-genres-dataset/pop/pop.00027.wav'
```

### 1.0.8 Play audio

```
[6]: ipd.Audio(music_data[SAMPLE_FILE]["filepath"])
```

```
[6]: <IPython.lib.display.Audio object>
```

### 1.0.9 Duration

```
[7]: for filepath in music_data.keys():
      music_data[filepath]["duration"] = librosa.get_duration(filename=filepath)

# Show for a sample file
music_data[SAMPLE_FILE]["duration"]
```

```
[7]: 30.00018140589569
```

## 1.1 What does an audio signal comprise?

It is a three-dimensional signal in which three axes represent time, amplitude and frequency.


### 1.1.1 Waveform

A Waveplot lets us know the loudness of the audio at a given time.

Note: The recommended sampling rate and also the default is **22kHz**. We are not overriding it.

```
[8]: for filepath in music_data.keys():
      music_data[filepath]["waveform"] = music_data[filepath]["waveform"],
      ↪music_data[filepath]["sampling_rate"] = librosa.load(filepath)

# Show for a sample file
plt.figure(figsize=(16, 2))
librosa.display.waveshow(music_data[SAMPLE_FILE]['waveform'],
      ↪sr=music_data[SAMPLE_FILE]['sampling_rate'], alpha=1)
plt.show()
```



### 1.1.2 Beats, Tempo, and Predominant local pulse (PLP)

- **Beat**: is the regularly occurring pattern of rhythmic stresses in music.
- **Tempo**: is the speed of the Beat, usually expressed in Beats Per Minute (BPM)
- **Predominant local pulse (PLP)** analyzes the frequency domain to find a locally stable tempo for each frame.

```
[9]: for filepath in music_data.keys():
    music_data[filepath]["beats"] = librosa.beat.
    ↳beat_track(y=music_data[filepath]['waveform'],
    ↳sr=music_data[filepath]['sampling_rate'])[0]
    music_data[filepath]["tempo"] = librosa.beat.
    ↳tempo(y=music_data[filepath]['waveform'],
    ↳sr=music_data[filepath]['sampling_rate'])[0]
    music_data[filepath]["plp"] = librosa.beat.
    ↳plp(y=music_data[filepath]['waveform'],
    ↳sr=music_data[filepath]['sampling_rate'])[0]

# Show for a sample file
print(f"Beats: {music_data[SAMPLE_FILE]['beats']}")
print(f"Tempo (Beats per minute): {music_data[SAMPLE_FILE]['tempo']}")
print(f"Predominant Local Pulse (PLP): {music_data[SAMPLE_FILE]['plp']}")
```

Beats: 83.35433467741936

Tempo (Beats per minute): 83.35433467741936

Predominant Local Pulse (PLP): 0.0

### 1.1.3 Spectrogram

- A spectrogram shows different frequencies playing at a particular time along with it's amplitude.
- It does so by converting the audio data into short term Fourier transform!

```
[10]: for filepath in music_data.keys():
    X = librosa.stft(music_data[filepath]['waveform'])
    Xdb = librosa.amplitude_to_db(abs(X))
    music_data[filepath]["spectrogram"] = np.mean(Xdb)
```


```

# Show for a sample file
print(f'Mean of Spectrogram: {music_data[SAMPLE_FILE] ["spectrogram"]}')

# Display Spectrogram
X = librosa.stft(music_data[SAMPLE_FILE] ['waveform'])
Xdb = librosa.amplitude_to_db(abs(X))
plt.figure(figsize=(18, 5))
librosa.display.specshow(Xdb, sr=music_data[SAMPLE_FILE] ["sampling_rate"],
    ↪x_axis='time', y_axis='hz')
plt.colorbar()
plt.show()

```

Mean of Spectrogram: -4.269062519073486



#### 1.1.4 Chromagram

- It represents the intensity of the twelve distinctive pitch classes that are used to study music.
- They can be employed in the differentiation of the pitch class profiles between audio signals.

```

[11]: for filepath in music_data.keys():
    music_data[filepath] ["chromagram"] = np.mean(librosa.feature.
    ↪chroma_stft(y=music_data[filepath] ['waveform'],
    ↪sr=music_data[filepath] ['sampling_rate']))


# Show for a sample file
print(f'Mean of Chromagram: {music_data[SAMPLE_FILE] ["chromagram"]}')

# Display Chromagram
X = librosa.feature.chroma_stft(y=music_data[SAMPLE_FILE] ['waveform'],
    ↪sr=music_data[SAMPLE_FILE] ['sampling_rate'])
plt.figure(figsize=(18, 5))
librosa.display.specshow(X, sr=music_data[SAMPLE_FILE] ["sampling_rate"],
    ↪x_axis='time', y_axis='hz')

```

```
plt.colorbar()
plt.show()
```

Mean of Chromagram: 0.3776119649410248



### 1.1.5 Tempogram


- Tempo can vary locally within a piece.
- Therefore, the tempogram is a feature matrix which indicates the prevalence of certain tempi at each moment in time.

```
[12]: for filepath in music_data.keys():
    hop_length = 512
    oenv = librosa.onset.onset_strength(y=music_data[filepath]['waveform'],
    ↪sr=music_data[filepath]['sampling_rate'], hop_length=hop_length)
    music_data[filepath]["tempogram"] = np.mean(librosa.feature.
    ↪tempogram(onset_envelope=oenv, sr=music_data[filepath]['sampling_rate'],
    ↪hop_length=hop_length))

# Show for a sample file
print(f'Mean of Tempogram: {music_data[SAMPLE_FILE]["tempogram"]}')

# Display Tempogram
hop_length = 512
oenv = librosa.onset.onset_strength(y=music_data[SAMPLE_FILE]['waveform'],
    ↪sr=music_data[SAMPLE_FILE]['sampling_rate'], hop_length=hop_length)
X = librosa.feature.tempogram(onset_envelope=oenv,
    ↪sr=music_data[SAMPLE_FILE]['sampling_rate'], hop_length=hop_length)
plt.figure(figsize=(18, 5))
librosa.display.specshow(X, sr=music_data[SAMPLE_FILE]["sampling_rate"],
    ↪x_axis='time', y_axis='hz')
plt.colorbar()
plt.show()
```

Mean of Tempogram: 0.1270049797896095



### 1.1.6 Zero Crossings and Zero Crossing Rate

- The zero crossing rate is the rate of sign-changes along a signal, i.e., the rate at which the signal changes from positive to negative or back.
- This feature has been used heavily in both speech recognition and music information retrieval.
- It usually has higher values for highly percussive sounds like those in metal and rock.

```
[13]: for filepath in music_data.keys():
    music_data[filepath]["zero_crossings"] = sum(librosa.
    ↪zero_crossings(music_data[filepath]['waveform'], pad=False))
    music_data[filepath]["zero_crossing_rate"] = np.mean(librosa.feature.
    ↪zero_crossing_rate(music_data[filepath]['waveform'], pad=False))

# Show for a sample file
print(f'Number of Zero crossings: {music_data[SAMPLE_FILE]["zero_crossings"]}')
print(f'Mean of Zero crossing rate: ␣
    ↪{music_data[SAMPLE_FILE]["zero_crossing_rate"]}')

```

Number of Zero crossings: 67487

Mean of Zero crossing rate: 0.1018672509425754

### 1.1.7 Spectral Centroid

- It indicates where the “centre of mass” for a sound is located.
- It is calculated as the weighted mean of the frequencies present in the sound.
- If the frequencies in music are same throughout then spectral centroid would be around a centre and if there are high frequencies at the end of sound then the centroid would be towards its end.

```
[14]: for filepath in music_data.keys():

```

```

spectral_centroids = librosa.feature.
↳spectral_centroid(y=music_data[filepath]['waveform'],
↳sr=music_data[filepath]['sampling_rate'])[0]
music_data[filepath]["spectral_centroid"] = np.mean(spectral_centroids)

# Computing the time variable for visualization
frames = range(len(spectral_centroids))
music_data[filepath]["time_variable"] = librosa.frames_to_time(frames)

# Show for a sample file
print(f'Mean of Spectral Centroid:
↳{music_data[SAMPLE_FILE]["spectral_centroid"]}')

# Computing the time variable for visualization
spectral_centroids = librosa.feature.
↳spectral_centroid(y=music_data[SAMPLE_FILE]['waveform'],
↳sr=music_data[SAMPLE_FILE]['sampling_rate'])[0]
frames = range(len(spectral_centroids))
t = librosa.frames_to_time(frames)


# Normalising the spectral centroid for visualisation
def normalize(x, axis=0):
    return sklearn.preprocessing.minmax_scale(x, axis=axis)

# Plotting the Spectral Centroid along the waveform
plt.figure(figsize=(18, 5))
librosa.display.waveshow(music_data[SAMPLE_FILE]['waveform'],
↳sr=music_data[SAMPLE_FILE]['sampling_rate'], alpha=0.4)
plt.plot(t, normalize(spectral_centroids), color='r')

```

Mean of Spectral Centroid: 2682.2841411936142

[14]: [<matplotlib.lines.Line2D at 0x7f9ca7c1d6a0>]





### 1.1.8 Spectral Rolloff

- It is the per-frame center frequency for a spectrogram bin such that at least roll\_percent (0.85 by default) of the energy of the spectrum in this frame is contained in this bin and the bins below.
- This can be used to, e.g., approximate the maximum (or minimum) frequency by setting roll\_percent to a value close to 1 (or 0).


```
[15]: for filepath in music_data.keys():
    music_data[filepath]["spectral_rolloff"] = np.mean(librosa.feature.
    ↳spectral_rolloff(y=music_data[filepath]['waveform'],
    ↳sr=music_data[filepath]['sampling_rate'])[0])

# Show for a sample file
print(f'Mean of Spectral Rolloff:
    ↳{music_data[SAMPLE_FILE]["spectral_rolloff"]}')

# Plotting the Spectral Rolloff along the waveform
spectral_rolloff = librosa.feature.
    ↳spectral_rolloff(y=music_data[SAMPLE_FILE]['waveform'],
    ↳sr=music_data[SAMPLE_FILE]['sampling_rate'])[0]
plt.figure(figsize=(18, 5))
librosa.display.waveshow(music_data[SAMPLE_FILE]['waveform'],
    ↳sr=music_data[SAMPLE_FILE]['sampling_rate'], alpha=0.4)
plt.plot(t, normalize(spectral_rolloff), color='r')
```

Mean of Spectral Rolloff: 5956.0789882721865

```
[15]: [<matplotlib.lines.Line2D at 0x7f9caa55a7c0>]
```



### 1.1.9 Spectral Contrast

- Each frame of a spectrogram  $S$  is first divided into sub-bands.

- Then, for each sub-band, the energy contrast is estimated by comparing the mean energy in the top quantile (peak energy) to that of the bottom quantile (valley energy).
- High contrast values generally correspond to clear, narrow-band signals, while low contrast values correspond to broad-band noise.


```
[16]: for filepath in music_data.keys():
    music_data[filepath]["spectral_contrast"] = np.mean(librosa.feature.
    ↪spectral_contrast(y=music_data[filepath]['waveform'],
    ↪sr=music_data[filepath]['sampling_rate'])[0])

# Show for a sample file
print(f'Mean of Spectral Contrast:
    ↪{music_data[SAMPLE_FILE]["spectral_contrast"]}')

# Plotting the Spectral Contrast along the waveform
spectral_contrast = librosa.feature.
    ↪spectral_contrast(y=music_data[SAMPLE_FILE]['waveform'],
    ↪sr=music_data[SAMPLE_FILE]['sampling_rate'])[0]
plt.figure(figsize=(18, 5))
librosa.display.waveshow(music_data[SAMPLE_FILE]['waveform'],
    ↪sr=music_data[SAMPLE_FILE]['sampling_rate'], alpha=0.4)
plt.plot(t, normalize(spectral_contrast), color='r')
```

Mean of Spectral Contrast: 22.257477751200955

[16]: [<matplotlib.lines.Line2D at 0x7f9cab0ffb20>]



### 1.1.10 Spectral Bandwidth

- Also called the spectral spread.
- It is the spectral range of interest around the centroid, that is, the variance from the spectral centroid.
- It is directly correlated with the perceived timbre.


```
[17]: for filepath in music_data.keys():
        music_data[filepath]["spectral_bandwidth"] = np.mean(librosa.feature.
        ↪spectral_bandwidth(y=music_data[filepath]['waveform'],
        ↪sr=music_data[filepath]['sampling_rate']))[0])

# Show for a sample file
print(f'Mean of Spectral Bandwidth:
    ↪{music_data[SAMPLE_FILE]["spectral_bandwidth"]}')

# Plotting the Spectral Bandwidth along the waveform
spectral_bandwidth = librosa.feature.
    ↪spectral_bandwidth(y=music_data[SAMPLE_FILE]['waveform'],
    ↪sr=music_data[SAMPLE_FILE]['sampling_rate'])[0]
plt.figure(figsize=(18, 5))
librosa.display.waveshow(music_data[SAMPLE_FILE]['waveform'],
    ↪sr=music_data[SAMPLE_FILE]['sampling_rate'], alpha=0.4)
plt.plot(t, normalize(spectral_bandwidth), color='r')
```

Mean of Spectral Bandwidth: 2923.9677291945763

```
[17]: [<matplotlib.lines.Line2D at 0x7f9cb69c8eb0>]
```



### 1.1.11 Spectral Flatness

- Also called the Tonality Coefficient.
- It is a measure to quantify how much noise-like a sound is, as opposed to being tone-like.
- A high spectral flatness (closer to 1.0) indicates the spectrum is similar to white noise. It is often converted to decibel.

```
[18]: for filepath in music_data.keys():
        music_data[filepath]["spectral_flatness"] = np.mean(librosa.feature.
        ↪spectral_flatness(y=music_data[filepath]['waveform']))[0])

# Show for a sample file
```

```


print(f'Mean of Spectral Flatness:␣
↪{music_data[SAMPLE_FILE] ["spectral_flatness"]}')

# Plotting the Spectral Flatness along the waveform
spectral_flatness = librosa.feature.
↪spectral_flatness(y=music_data[SAMPLE_FILE] ['waveform']) [0]
plt.figure(figsize=(18, 5))
librosa.display.waveshow(music_data[SAMPLE_FILE] ['waveform'],␣
↪sr=music_data[SAMPLE_FILE] ['sampling_rate'], alpha=0.4)
plt.plot(t, normalize(spectral_flatness), color='r')

```

Mean of Spectral Flatness: 0.047302521765232086

[18]: [matplotlib.lines.Line2D at 0x7f9c9923a0d0]



### 1.1.12 RMS (Root Mean Square) Energy

- Compute root-mean-square (RMS) energy for each frame from the audio samples.
- For audio signals, that roughly corresponds to how loud the signal is.

```

[19]: for filepath in music_data.keys():
    music_data[filepath] ["rms"] = np.mean(librosa.feature.
↪rms(y=music_data[filepath] ['waveform']) [0])


# Show for a sample file
print(f'Mean of RMS (Root Mean Square) Energy:␣
↪{music_data[SAMPLE_FILE] ["rms"]}')

# Plotting the Spectral Flatness along the waveform
rms = librosa.feature.rms(y=music_data[SAMPLE_FILE] ['waveform']) [0]
plt.figure(figsize=(18, 5))
librosa.display.waveshow(music_data[SAMPLE_FILE] ['waveform'],␣
↪sr=music_data[SAMPLE_FILE] ['sampling_rate'], alpha=0.4)
plt.plot(t, normalize(rms), color='r')

```

Mean of RMS (Root Mean Square) Energy: 0.1579022854566574

[19]: [



### 1.1.13 MFCC — Mel-Frequency Cepstral Coefficients


- One of the most important method to extract a feature of an audio signal.
- They are a small set of features (usually about 10–20) which concisely describe the overall shape of a spectral envelope.

```
[20]: for filepath in music_data.keys():
    music_data[filepath]["mfccs"] = np.mean(librosa.feature.
    ↳mfcc(y=music_data[filepath]['waveform'],
    ↳sr=music_data[filepath]['sampling_rate']))

# Show for a sample file
print(f'Mean of MFCC - Mel-Frequency Cepstral Coefficients:
    ↳{music_data[SAMPLE_FILE]["mfccs"]}')

# Displaying the MFCCs:
mfccs = librosa.feature.mfcc(y=music_data[SAMPLE_FILE]['waveform'],
    ↳sr=music_data[SAMPLE_FILE]['sampling_rate'])
plt.figure(figsize=(18, 5))
librosa.display.specshow(mfccs, sr=music_data[SAMPLE_FILE]['sampling_rate'],
    ↳x_axis='time')
plt.show()
```

Mean of MFCC - Mel-Frequency Cepstral Coefficients: 2.0804495811462402



### 1.1.14 Prepare Feature Dataset

```
[21]: df_feature = pd.DataFrame(music_data.values())
df_feature_with_filenames = df_feature.
      drop(columns=['waveform', 'time_variable', 'sampling_rate', 'duration'])
df_feature_with_filenames.head()
```

```
[21]:
```

	filepath	filename	genre	\
0	data/music-genres-dataset/pop/pop.00027.wav	pop.00027.wav	pop	
1	data/music-genres-dataset/pop/pop.00033.wav	pop.00033.wav	pop	
2	data/music-genres-dataset/pop/pop.00032.wav	pop.00032.wav	pop	
3	data/music-genres-dataset/pop/pop.00026.wav	pop.00026.wav	pop	
4	data/music-genres-dataset/pop/pop.00030.wav	pop.00030.wav	pop	

	beats	tempo	plp	spectrogram	chromagram	tempogram	\
0	83.354335	83.354335	0.000000	-4.269063	0.377612	0.127005	
1	99.384014	99.384014	0.000000	-3.781867	0.423036	0.158264	
2	95.703125	95.703125	0.519045	-4.358516	0.426769	0.121314	
3	75.999540	151.999081	0.000000	-4.899115	0.334962	0.133198	
4	95.703125	95.703125	0.472598	0.983008	0.357630	0.140248	

	zero_crossings	zero_crossing_rate	spectral_centroid	spectral_rolloff	\
0	67487	0.101867	2682.284141	5956.078988	
1	90785	0.137010	2835.119375	5944.979313	
2	151492	0.228641	4225.461049	8483.382577	
3	68571	0.103494	2525.591959	5552.801896	
4	98516	0.148679	3070.672741	6653.076965	

	spectral_contrast	spectral_bandwidth	spectral_flatness	rms	\
0	22.257478	2923.967729	0.047303	0.157902	
1	18.068245	2731.752387	0.061382	0.139896	
2	18.657617	3404.380733	0.125869	0.097861	
3	23.054193	2848.409561	0.030347	0.195626	

4	17.128994	2975.765677	0.075034	0.203093
---	-----------	-------------	----------	----------

```
mfccs
0 2.080450
1 0.843129
2 -0.666234
3 2.070004
4 2.515382
```

### 1.1.15 Prepare Training Dataset: Drop irrelevant columns

```
[22]: df_training_data = df_feature.
      ↪drop(columns=['waveform','time_variable','filepath','sampling_rate','duration',
      ↪'filename'])
      df_training_data.head()
```

```
[22]: genre      beats      tempo      plp  spectrogram  chromagram  tempogram  \
0  pop  83.354335  83.354335  0.000000   -4.269063    0.377612    0.127005
1  pop  99.384014  99.384014  0.000000   -3.781867    0.423036    0.158264
2  pop  95.703125  95.703125  0.519045   -4.358516    0.426769    0.121314
3  pop  75.999540  151.999081  0.000000   -4.899115    0.334962    0.133198
4  pop  95.703125  95.703125  0.472598    0.983008    0.357630    0.140248
```

	zero_crossings	zero_crossing_rate	spectral_centroid	spectral_rolloff	\
0	67487	0.101867	2682.284141	5956.078988	
1	90785	0.137010	2835.119375	5944.979313	
2	151492	0.228641	4225.461049	8483.382577	
3	68571	0.103494	2525.591959	5552.801896	
4	98516	0.148679	3070.672741	6653.076965	

	spectral_contrast	spectral_bandwidth	spectral_flatness	rms	\
0	22.257478	2923.967729	0.047303	0.157902	
1	18.068245	2731.752387	0.061382	0.139896	
2	18.657617	3404.380733	0.125869	0.097861	
3	23.054193	2848.409561	0.030347	0.195626	
4	17.128994	2975.765677	0.075034	0.203093	

```
mfccs
0 2.080450
1 0.843129
2 -0.666234
3 2.070004
4 2.515382
```

### 1.1.16 Save both datasets

```
[23]: if config["persist"]:
        df_feature_with_filenames.to_csv("data/" + config["dataset"]["output"] +
        ↪ "-with-filenames.csv", sep="\t", header=True, index=True)
        df_training_data.to_csv("data/" + config["dataset"]["output"] + ".csv",
        ↪ sep="\t", header=True, index=True)
```